

## BOSC2013 OO2 - Asynchronicity and C

Sigurt Bladt Dinesen  
sidi@itu.dk

November 18, 2013

## Contents

<b>1</b>	<b>Multi-threaded sum</b>	<b>2</b>
1.1	A faster sum of $n$ square-roots . . . . .	2
1.2	Testing . . . . .	3
<b>2</b>	<b>Multi-threaded FIFO buffer as a linked list</b>	<b>3</b>
2.1	Parallelized read/write actions . . . . .	3
2.2	Thread Safety . . . . .	4
<b>3</b>	<b>Producer-Consumer with a bounded buffer</b>	<b>4</b>
3.1	Termination . . . . .	4
3.2	Testing . . . . .	4
<b>4</b>	<b>Banker's algorithm</b>	<b>5</b>
4.1	C Matrix Allocation . . . . .	5
4.2	State-safety . . . . .	5
<b>5</b>	<b>Appendix/code</b>	<b>6</b>
5.1	Multi-threaded sum . . . . .	6
5.2	Multi-threaded FIFO buffer as a linked list . . . . .	7
5.3	Producer-Consumer with a bounded buffer . . . . .	12
5.4	Banker's Algorithm . . . . .	16

# 1 Multi-threaded sum

## 1.1 A faster sum of $n$ square-roots

*Rewrite the program in such a way that it will actually run faster, than it would have on a single thread.*

This is achieved by letting  $m$  threads sum the squares of  $\frac{n}{m}$  numbers each in parallel, and then summing the results in linear time afterwards. The advantage of this approach is that the use of mutexes is redundant: Each thread keeps its own tentative sum, the collection of which is summed only after all threads finish.

A more elegant approach would be to run the tentative sums, and then have each thread end by locking a mutex, adding its sum to a common sum variable, and unlocking the mutex. The later solution would also offer a performance increase proportional to the ratio  $\frac{n}{m}$ , which will usually be large.

To conclude the assignment: The runtime of *sumsqrt* is now  $\frac{n}{m} + m$  rather than  $n$  as it was initially (provided at least  $m$  logical cores of course).

A visual representation of the runtime of *sumsqrt* can be seen in figure 1. It can be seen that on one core, the real time and cpu time spend in user land are equal, on two cores the real time is almost halved, although the user land cpu time increases slightly, probably due to threading overhead. It can also be seen that threads beyond the number of cores does virtually nothing to either times. Kernel mode cpu time is omitted as it is negligible here. Had the program used mutexes, that may not have been the case.

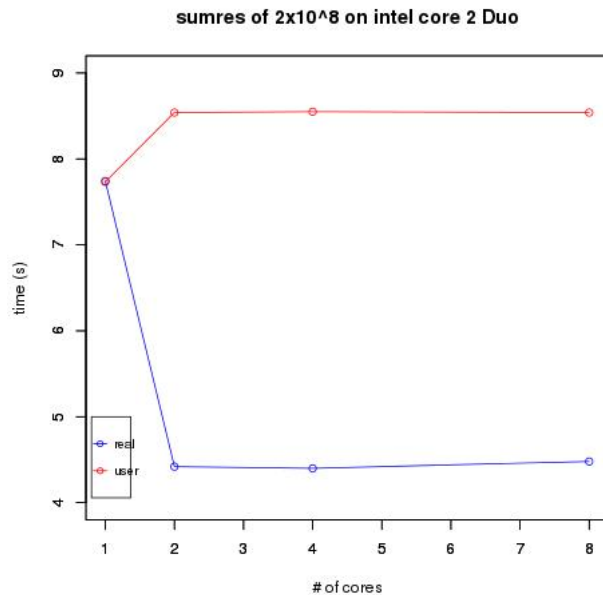


Figure 1: Running *sumsqrt* on the first 200 million natural numbers on 1, 2, 4, and 8 threads on an 1.6 Ghz Intel core 2 Duo processor. Data was collected over a total of 12 runs.

## 1.2 Testing

To test *sumsqrt*, first compile by running

```
make sumsqrt
```

in the project root directory. Then *sumsqrt* can be used as

```
sumsqrt [numbers to root and sum] [number of threads]
```

For example

```
sumsqrt 100 4
```

Will calculate the sum of the square roots of numbers 0 though 99, using 4 threads.

## 2 Multi-threaded FIFO buffer as a linked list

### 2.1 Parallelized read/write actions

*Assume that the list implemetation is used in an asynchronous environment. What issues might occur?*

Since *list\_remove* and *list\_add* work on the same list (albeit on different ends of it), raise conditions would occur if different threads of execution were to run these methods simultaneously. In fact, different threads running just *one* of these methods would be problematic. The trivial example is that of reading from, and writing to, a variable, e.g. to increment its value by a constant.

The execution of such example is illustrated in figure 1, where  $T_2$  sets the shared entity to the result applying a function to it, regardless of changes made by  $T_1$ . After execution,  $S$  is the result of applying  $T_2.f$  to  $\alpha$ , when it should have been the result of applying  $T_2.f$  to  $\alpha'$ . The thing to notice here is that

$$T_2.f(\alpha) \neq T_2.f(\alpha') = T_2.f(T_1.f(\alpha))$$

$T_1$ action	$T_2$ action	$S$ value
		$\alpha$
$T_1.a = S$		$\alpha$
	$T_2.a = S$	$\alpha$
$S = T_1.f(S) = \alpha'$		$T_1.f(\alpha)$
	$S = T_2.f(S)$	$T_2.f(\alpha)$

Table 1: Execution of two threads on a shared entity. The figure shows the actions of two threads,  $T_1$  and  $T_2$ , operating on a shared value  $S$ .

Similar problems would occur for the next variables of the list's nodes.

## 2.2 Thread Safety

*Use a mutex to make a thread-safe version of list.c*

This is achieved quite simply by using the C pthread library to take a lock before either adding or removing elements in the buffer, and releasing it when done. Packaging all write operations on the list into a critical section ensures that the situation in table 1 can never occur. Note that *add* and *remove* uses the same mutex. This is necessary as the two functions sometimes will operate on the same variables.

## 3 Producer-Consumer with a bounded buffer

The producer-controller is implemented with the thread safe list implementation from last section. This is not enough however, as we wish to ensure *list\_add* and *list\_remove* operations will wait if the list is full or empty, respectively. This demand is met by the use of two semaphores, one for adding (*empty*) and one for removing (*full*). When a producer wishes to add a product, preforms *sem\_wait* on *empty*, to ensure there are empty slots in the buffer, and then performs a *sem\_post* on *full* to indicate consumers that there is one more full slot. Consumer threads do the opposite, that is *sem\_post* is performed on *empty* and *sem\_wait* on *full*. The *full* semaphore is initialized to 0, and *empty* to the desired buffer size. This corresponds to the notion that *full* represents occupied buffer slots (consumable slots), and *empty* represent empty slots (available to the producers).

### 3.1 Termination

*Make sure the program does not terminate before all products are produced and consumed*

This is done by keeping the thread ids for consumer and producer threads in two lists, and joining on them all in the end of the programs main method. This guarantees that program terminates only when all consumers and producers are done. The number of product consumed or produced by each thread is determined at runtime to be the ratio between the number of products and threads. This ensures that all produced products will be consumed, regardless of whether the number of producers and consumers are equal.

### 3.2 Testing

To test the producer-consumer program (and with it the FIFO buffer) compile by running

```
make prodcons
```

in the project root directory. Then run prodcons as

```
prodcons [number of prodcer threads] [number of consumer threads]  
[buffer size] [number of products]
```

For example,

```
./prodcons 10 10 10 20
```

Will run the producer-consumer program with 10 producing threads, each producing 2 products, 10 consuming threads, each consuming 2 products, a buffer with 10 slots, producing a total of 20 products.

## 4 Banker's algorithm

### 4.1 C Matrix Allocation

*Allocate memory for the state's vectors and matrices dynamically*

The rows of each matrix are allocated with *malloc*, as C does not have a syntax for allocating nested arrays. This memory is never explicitly freed, as the arrays are used until program termination anyway.

### 4.2 State-safety

*Use the Banker's safety algorithm to determine whether a state is safe*

The function *resource\_request* implements both the request and safety algorithm. It takes an array as argument, representing a resource allocation request, and then runs the safety algorithm as if the request had already been granted.

The *resource\_request* and *resource\_release* functions both operate on the same matrices, and therefore lock the same mutex before writing to them to avoid race conditions. Using the same mutex ensures that deadlocks can never occur, as a deadlock needs a minimum of two locks to wait on.

Before starting the requesting threads, the program runs the safety algorithm with a request of 0 of all resources, thus ensuring that the current state (equivalent to the state that results of requesting nothing) is safe.

## 5 Appendix/code

### 5.1 Multi-threaded sum

sumsqrt.c

```
1 #include <pthread.h>
2 #include <stdio.h>
3 #include <math.h>
4
5 typedef struct runpars{
6     int lower;
7     int upper;
8     float res;
9 } runpars;
10
11 void *runner(void *params);
12 void runThread(pthread_t *tidP, runpars *params);
13 void spawnNThreads(int n, runpars pars[]);
14
15 int main(int argc, char* argv[]){
16     if(argc != 3){
17         fprintf(stderr, "usage: sumsqrt [n] [
18             threads]\n");
19         return -1;
20     }
21     if(atoi(argv[1]) < 0) {
22         fprintf(stderr, "%d - n must be >= 0\n",
23             atoi(argv[1]));
24         return -1;
25     }
26     if(atoi(argv[2]) <= 0) {
27         fprintf(stderr, "%d - threads must be = 0\
28             n", atoi(argv[2]));
29         return -1;
30     }
31
32     int upper = atoi(argv[1]);
33     int nThreads = atoi(argv[2]);
34     int stepSize = upper/nThreads;
35     struct runpars pars[nThreads];
36     int i;
37     for (i=0; i<nThreads; i++){
38         int first = i * stepSize;
39         int last = (i == nThreads-1) ? upper : (i
40             +1) * stepSize -1;
41         pars[i] = (runpars) {.lower = first, .
42             upper = last, .res = 0.0};
43         //printf("%d to %d\n", pars[i].lower,
44             pars[i].upper);
45     }
```

```

40     spawnNThreads(nThreads, pars);
41     float sum = 0.0;
42     for (i=0; i<nThreads; i++){
43         sum += pars[i].res;
44     }
45     printf("sum of squares 0 though %d = %f\n", upper
        , sum);
46
47     return 0;
48 }
49
50 void spawnNThreads(int n, runpars pars[]) {
51     pthread_t ids[n];
52     int i;
53     for (i=0; i<n; i++){
54         runThread(&(ids[i]), &(pars[i]));
55     }
56     for (i=0; i<n; i++)
57         pthread_join(ids[i], NULL);
58 }
59
60 void runThread(pthread_t *tidP, runpars *params){
61     pthread_attr_t attrP;
62     pthread_attr_init(&attrP);
63     pthread_create(tidP, &attrP, runner, params);
64     pthread_attr_destroy(&attrP);
65 }
66
67 void *runner(void *params){
68     struct runpars* pars = (struct runpars*) params;
69     int i, upper = pars->upper;
70     for(i = pars->lower; i <= upper; i++){
71         pars->res += sqrtf(i);
72     }
73     pthread_exit(0);
74 }

```

## 5.2 Multi-threaded FIFO buffer as a linked list

list/main.c

```

1  /*****
2
3      main.c
4
5      Implementation of a simple FIFO buffer as a
6      linked list defined in list.h.
7
8  *****/

```



```

7
8 #include <stdio.h>
9 #include <stdlib.h>
10 #include <pthread.h>
11 #include "list.h"
12 #include "assert.h"
13
14 // FIFO list;
15 List *fifo;
16 const int elmsPerThread = 500;
17 const int numberOfThreads = 3;
18
19 void *asyncRunner(void* peh){
20     int i;
21     for(i=0; i<elmsPerThread; i++){
22         list_add(fifo, node_new());
23     }
24     pthread_exit(NULL);
25 }
26 void asyncListTest(){ //this test is fake and gay, and
    not unit-like at all
27     fifo = list_new();
28     pthread_t *ids[numberOfThreads];
29     int i;
30     for(i=0; i<numberOfThreads; i++){
31         pthread_t* tid = (pthread_t *) malloc(
            sizeof(pthread_t));
32         ids[i] = tid;
33         pthread_create(tid, NULL, asyncRunner,
            NULL);
34     }
35     for(i=0; i<numberOfThreads; i++){
36         pthread_join(*ids[i], NULL);
37     }
38     int removeN=0;
39     int actualN=0;
40     int assertedN = fifo->len;
41     Node *n = fifo->first->next;
42     while(n != NULL) {
43         actualN++;
44         n = n->next;
45     }
46
47     n = list_remove(fifo);
48     while(n != NULL){
49         removeN++;
50         n = list_remove(fifo);
51     }
52     //printf("asserted: %i\nremoved: %i\nactual: %i\n
        elmsPer: %i\n#of threads: %i\nprod: %i\n",

```

```

        assertedN, removeN, actualN, elmsPerThread,
        numberOfThreads, elmsPerThread*numberOfThreads
    );
53  assert(assertedN == actualN);
        //fifo ->len is the same
        as the actual length
54  assert(removeN == actualN);
        //remove will yield as
        many elms as there actually are in the list
55  assert(elmsPerThread*numberOfThreads == actualN);
        //there are actually the expected number of
        elements
56  assert(fifo->len == 0);
        //the list is now
        empty
57 }
58
59 int main(int argc, char* argv[])
60 {
61     asyncListTest();
62     return 0;
63     fifo = list_new();
64
65     list_add(fifo, node_new_str("s1"));
66     list_add(fifo, node_new_str("s2"));
67
68     Node *n1 = list_remove(fifo);
69     if (n1 == NULL) { printf("Error no elements in
        list\n"); exit(-1);}
70     Node *n2 = list_remove(fifo);
71     if (n2 == NULL) { printf("Error no elements in
        list\n"); exit(-1);}
72     printf("%s\n%s\n", n1->elm, n2->elm);
73
74     return 0;
75 }

```

list/list.h

```

1  /*****
2
3      list.h
4
5      Header file with definition of a simple linked list.
6
7      *****/
8  #ifndef _LIST_H
9  #define _LIST_H

```

```

10
11 /* structures */
12 typedef struct node {
13     void *elm; /* use void type for generality; we cast the
14                element's type to void type */
15     struct node *next;
16 } Node;
17
18 typedef struct list {
19     int len;
20     Node *first;
21     Node *last;
22 } List;
23
24 /* functions */
25 List *list_new(void); /* return a new list
26                        structure */
27 void list_add(List *l, Node *n); /* add node n to list l
28                                as the last element */
29 Node *list_remove(List *l); /* remove and return the
30                             first element from list l*/
31 Node *node_new(void); /* return a new node
32                       structure */
33 Node *node_new_str(char *s); /* return a new node
34                              structure, where elm points to new copy of string s */
35 #endif

```

list/main.c

```

1  /*****
2
3      main.c
4
5      Implementation of a simple FIFO buffer as a
6      linked list defined in list.h.
7
8  *****/
9
10 #include <stdio.h>
11 #include <stdlib.h>
12 #include <pthread.h>
13 #include "list.h"
14 #include "assert.h"
15
16 // FIFO list;
17 List *fifo;
18 const int elmsPerThread = 500;
19 const int numberOfThreads = 3;

```

```

18
19 void *asyncRunner(void* peh){
20     int i;
21     for(i=0; i<elmsPerThread; i++){
22         list_add(fifo, node_new());
23     }
24     pthread_exit(NULL);
25 }
26 void asyncListTest(){ //this test is fake and gay, and
    not unit-like at all
27     fifo = list_new();
28     pthread_t *ids[numberOfThreads];
29     int i;
30     for(i=0; i<numberOfThreads; i++){
31         pthread_t* tid = (pthread_t *) malloc(
32             sizeof(pthread_t));
33         ids[i] = tid;
34         pthread_create(tid, NULL, asyncRunner,
35             NULL);
36     }
37     for(i=0; i<numberOfThreads; i++){
38         pthread_join(*ids[i], NULL);
39     }
40     int removeN=0;
41     int actualN=0;
42     int assertedN = fifo->len;
43     Node *n = fifo->first->next;
44     while(n != NULL) {
45         actualN++;
46         n = n->next;
47     }
48     n = list_remove(fifo);
49     while(n != NULL){
50         removeN++;
51         n = list_remove(fifo);
52     }
53     //printf("asserted: %i\nremoved: %i\nactual: %i\n
54         elmsPer: %i\n#of threads: %i\nprod: %i\n",
55         assertedN, removeN, actualN, elmsPerThread,
56         numberOfThreads, elmsPerThread*numberOfThreads
57     );
58     assert(assertedN == actualN);
59     //fifo->len is the same
60     as the actual length
61     assert(removeN == actualN);
62     //remove will yield as
63     many elms as there actually are in the list
64     assert(elmsPerThread*numberOfThreads == actualN);
65     //there are actually the expected number of

```

```

56         elements
           assert(fifo->len == 0);
                                           //the list is now
           empty
57     }
58
59 int main(int argc, char* argv[])
60 {
61     asyncListTest();
62     return 0;
63     fifo = list_new();
64
65     list_add(fifo, node_new_str("s1"));
66     list_add(fifo, node_new_str("s2"));
67
68     Node *n1 = list_remove(fifo);
69     if (n1 == NULL) { printf("Error no elements in
                           list\n"); exit(-1);}
70     Node *n2 = list_remove(fifo);
71     if (n2 == NULL) { printf("Error no elements in
                           list\n"); exit(-1);}
72     printf("%s\n%s\n", n1->elm, n2->elm);
73
74     return 0;
75 }

```

list/Makefile

```

1  all: fifo liblist.a
2
3  OBJS = list.o main.o
4  LIBS= -lpthread
5  CFLAGS= "-ggdb"
6
7  fifo: ${OBJS}
8         gcc -o $@ ${OBJS} ${LIBS}
9
10 liblist.a: list.o
11         ar rcs liblist.a list.o
12
13 clean:
14         rm -rf *o fifo

```

### 5.3 Producer-Consumer with a bounded buffer

prodcons.c

```

1  #include <pthread.h>
2  #include <semaphore.h>
3  #include <unistd.h>

```

```

4 #include <stdio.h>
5 #include <stdlib.h>
6 #include "list/list.h"
7
8 typedef struct runpars{
9     int start;
10    int end;
11 } runpars;
12
13 void *producer(void *params);
14 void *consumer(void *params);
15 void runThread(pthread_t *tidP, runpars *params);
16 pthread_t* spawnNThreads(int n, void>(*fun) (void *));
17 void sleeper(float waitMs);
18 int prodN;
19 int consN;
20 int bufferSize;
21 int productsN;
22 int producerLoad;
23 int consumerLoad;
24 float SLEEP_TIME_MS = 1.0;
25 List *buffer;
26 sem_t empty, full;
27
28 int main(int argc, char* argv[]) {
29     setbuf(stdout, NULL);
30     if(argc != 5){
31         fprintf(stderr, "usage: prodcons [
32             procucers] [consumers] [bufferSize] [
33             products]\n");
34         return -1;
35     }
36     if(atoi(argv[1]) <= 0) {
37         fprintf(stderr, "%d - producers must be >
38             0\n", atoi(argv[1]));
39         return -1;
40     }
41     if(atoi(argv[2]) <= 0) {
42         fprintf(stderr, "%d - consumers must be >
43             0\n", atoi(argv[2]));
44         return -1;
45     }
46     if(atoi(argv[3]) <= 0) {
47         fprintf(stderr, "%d - bufferSize must be >
48             0\n", atoi(argv[3]));
49         return -1;
50     }
51     if(atoi(argv[4]) <= 0) {
52         fprintf(stderr, "%d - products must be >
53             0\n", atoi(argv[4]));
54         return -1;
55     }
56 }

```

```

48         return -1;
49     }
50
51     buffer = list_new();
52     prodN = atoi(argv[1]);
53     consN = atoi(argv[2]);
54     bufferSize = atoi(argv[3]);
55     productsN = atoi(argv[4]);
56     producerLoad = productsN / prodN;
57     consumerLoad = productsN / consN;
58     printf("Starting! Here are the stats\n"
59           "prodN: %i\n"
60           "consN: %i\n"
61           "buffer: %i\n"
62           "products: %i\n"
63           "produderload: %i\n"
64           "consumerload: %i\n",
65           prodN, consN, bufferSize,
66           productsN, producerLoad,
67           consumerLoad);
68
69     sem_init(&full, 0, 0);
70     sem_init(&empty, 0, bufferSize);
71     srand(time(NULL));
72
73     pthread_t *consIds = spawnNThreads(consN,
74                                         consumer);
75     pthread_t *prodIds = spawnNThreads(prodN,
76                                         producer);
77     int i;
78     for(i=0; i<prodN; i++){
79         pthread_join(prodIds[i], NULL);
80     }
81     for(i=0; i<consN; i++){
82         pthread_join(consIds[i], NULL);
83     }
84     free(prodIds);
85     free(consIds);
86     return 0;
87 }
88
89 pthread_t* spawnNThreads(int n, void (*fun)(void *)){
90     pthread_t *ids = (pthread_t*) malloc(sizeof(
91         pthread_t)*n);
92     int i;
93     for (i=0; i<n; i++){
94         struct runpars *par = malloc(sizeof(
95             runpars));
96         if(fun==producer){
97             par->start = producerLoad*i;

```

```

92         par->end = i==n-1 ? productsN -1 : (
93             producerLoad * (i+1))-1;
94         printf("produce steps: %i - %i\n", par->
95             start, par->end);
96     }else{
97         par->start = consumerLoad*i;
98         par->end = i==n-1 ? productsN -1 : (
99             consumerLoad * (i+1))-1;
100         printf("consume steps: %i - %i\n", par->
101             start, par->end);
102     }
103     pthread_create(&ids[i], NULL, fun, par);
104 }
105 return ids;
106 }
107
108 void *producer(void *params){
109     struct runpars* pars = (struct runpars*) params;
110     int i;
111     for(i = pars->start; i < pars->end; i++){
112         Node *n = node_new();
113         n->elm = (void*) i;
114         sem_wait(&empty);
115         list_add(buffer, n);
116         sem_post(&full);
117         sleeper(SLEEP_TIME_MS);
118         int fullC, emptyC;
119         sem_getvalue(&full, &fullC);
120         sem_getvalue(&empty, &emptyC);
121         printf("producer %i made item %i, full: %
122             i, empty: %i\n", pars->start/
123             producerLoad, (int) n->elm, fullC,
124             emptyC);
125     }
126     free(pars);
127     pthread_exit(0);
128 }
129
130 void *consumer(void *params){
131     struct runpars* pars = (struct runpars*) params;
132     int i;
133     for(i = pars->start; i < pars->end; i++){
134         sem_wait(&full);
135         Node *n = list_remove(buffer);
136         sem_post(&empty);
137         sleeper(SLEEP_TIME_MS);
138         int fullC, emptyC;
139         sem_getvalue(&full, &fullC);
140         sem_getvalue(&empty, &emptyC);
141         printf("consumer ate item %i, full: %i,

```



```

135         empty: %i\n", (int) n->elm, fullC,
136         }
137         free(pars);
138         pthread_exit(0);
139     }
140 void sleeper(float waitMs){
141     //random value between 0 and waitMs
142     waitMs = (float)rand() * waitMs / (float)RANDMAX
143     ;
144     usleep((int)waitMs * 1000); //times 1000 for
        microseconds;
}

```

makefile

```

1 prodcons: prodcons.o
2     gcc -g -static -Wall -I./list -L./list -o $@
        prodcons.o -llist -lpthread
3
4 report: report/report.pdf
5 report/report.pdf: report/report.tex
6     pdflatex --output-directory report/tmp report/
        report.tex
7     pdflatex --output-directory report/tmp report/
        report.tex #horribad, use latexmk or latex-MK
        instead
8     mv report/tmp/report.pdf report/report.pdf
9
10 sumsqrt: sumsqrt.o
11     gcc -ggdb sumsqrt.o -o sumsqrt -lpthread -lm
12
13 sumsqrt.o: sumsqrt.c
14     gcc -ggdb -c sumsqrt.c
15
16 clean:
17     @rm prodcons prodcons.o sumsqrt sumsqrt.o report
        /tmp/* report/report.pdf 2>/dev/null

```

## 5.4 Banker's Algorithm

banker.c

```

1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <sys/time.h>
4 #include <pthread.h>
5
6 typedef struct state {

```

```

7   int *resource;
8   int *available;
9   int **max;
10  int **allocation;
11  int **need;
12 } State;
13
14 // Global variables
15 int m, n;
16 State *s = NULL;
17 pthread_mutex_t mutex;
18
19 // Mutex for access to state.
20 pthread_mutex_t state_mutex;
21
22 void printavail(){
23     printf("Avaialble: ");
24     int i;
25     for(i=0; i<n; i++){
26         printf("%i,", s->available[i]);
27     }
28     puts("]");
29 }
30
31 /* Random sleep function */
32 void Sleep(float wait_time_ms)
33 {
34     // add randomness
35     wait_time_ms = ((float)rand())*wait_time_ms / (float)
        RANDMAX;
36     usleep((int) (wait_time_ms * 1e3f)); // convert from ms
        to us
37 }
38
39 //return 1 if a is <= b for all a[i] b[i], else return 0
40 int arrayLE(int* a, int* b, int len)
41 {
42     int i;
43     for(i=0; i<len; i++){
44         if(a[i] > b[i]){
45             return 0;
46         }
47     }
48     return 1;
49 }
50
51 /* Allocate resources in request for process i, only if
    it
52 results in a safe state and return 1, else return 0 */
53 int resource_request(int pi, int *request)

```

```

54 {
55     if (!arrayLE(request, s->available, n))
56         return 0;
57
58     int work[n];
59     int finish[m];
60     int tmpNeed[n];
61     int i;
62     for (i=0; i<m; i++){
63         finish[i]=0;
64     }
65     for (i=0; i<n; i++){
66         work[i]=s->available[i] - request[i];
67     }
68     for (i=0; i<n; i++){
69         tmpNeed[i]=s->need[pi][i] - request[i];
70     }
71     int k=0;
72     while(1){
73         while (!(!finish[k] && arrayLE(tmpNeed, work, n)) && k
74             <m){
75             k++;
76         }
77         if (!(!finish[k] && arrayLE(tmpNeed, work, n))){
78             break;
79         }
80
81         for (i=0; i<n; i++){
82             work[i] = work[i] + s->allocation[pi][k] + request[
83                 i];
84         }
85         finish[k]=1;
86         k=1;
87         for (i=0; i<m; i++){
88             if (!finish[i])
89                 k=0;
90             break;
91         }
92         if (k){
93             pthread_mutex_lock(&mutex);
94             for (i=0; i<n; i++){
95                 s->allocation[pi][i] = s->allocation[pi][i] +
96                     request[i];
97                 s->available[i] = s->available[i] - request[i];
98                 s->need[pi][i] = s->need[pi][i] - request[i];
99             }
100             printavail();
101             pthread_mutex_unlock(&mutex);

```

```

101     }
102     return k;
103 }
104
105 /* Release the resources in request for process i */
106 void resource_release(int i, int *request)
107 {
108     int j;
109     pthread_mutex_lock(&mutex);
110     for (j=0; j<n; j++){
111         s->available[j] = s->available[j] + request[j];
112         s->allocation[i][j] = s->allocation[i][j] - request[j]
113         ];
114         s->need[i][j] = s->need[i][j] + request[j];
115     }
116     pthread_mutex_unlock(&mutex);
117 }
118
119 /* Generate a request vector */
120 void generate_request(int i, int *request)
121 {
122     int j, sum = 0;
123     while (!sum) {
124         for (j = 0; j < n; j++) {
125             request[j] = s->need[i][j] * ((double)rand()) / (
126                 double)RANDMAX;
127             sum += request[j];
128         }
129     }
130     printf("Process %d: Requesting resources.\n", i);
131 }
132
133 /* Generate a release vector */
134 void generate_release(int i, int *request)
135 {
136     int j, sum = 0;
137     while (!sum) {
138         for (j = 0; j < n; j++) {
139             request[j] = s->allocation[i][j] * ((double)rand())
140             / (double)RANDMAX;
141             sum += request[j];
142         }
143     }
144     printf("Process %d: Releasing resources.\n", i);
145 }
146
147 /* Threads starts here */
148 void *process_thread(void *param)
149 {
150     /* Process number */

```

```

148     int i = (int) (long) param, j;
149     /* Allocate request vector */
150     int *request = malloc(n*sizeof(int));
151     while (1) {
152         /* Generate request */
153         generate_request(i, request);
154         while (!resource_request(i, request)) {
155             /* Wait */
156             Sleep(100);
157         }
158         /* Generate release */
159         generate_release(i, request);
160         /* Release resources */
161         resource_release(i, request);
162         /* Wait */
163         Sleep(1000);
164     }
165     free(request);
166 }
167
168 int main(int argc, char* argv[])
169 {
170     pthread_mutex_init(&mutex, NULL);
171
172     /* Get size of current state as input */
173     int i, j;
174     printf("Number of processes: ");
175     scanf("%d", &m);
176     printf("Number of resources: ");
177     scanf("%d", &n);
178
179     int res[n];
180     int avail[n];
181     int **max[m];
182     int **alloc[m];
183     int **need[m];
184
185     for(i=0; i<m; i++){
186         max[i] = malloc(sizeof(int)*n);
187         alloc[i] = malloc(sizeof(int)*n);
188         need[i] = malloc(sizeof(int)*n);
189     }
190
191     State p = (State)
192     {
193         .resource = (int*) res,
194         .available = (int*) avail,
195         .max = (int**) max,
196         .allocation = (int**) alloc,
197         .need = (int**) need

```

```

198     };
199     s = &p;
200     /* Get current state as input */
201     printf("Resource vector: ");
202     for(i = 0; i < n; i++)
203         scanf("%d", &s->resource[i]);
204     printf("Enter max matrix: ");
205     for(i = 0; i < m; i++)
206         for(j = 0; j < n; j++){
207             scanf("%d", &s->max[i][j]);
208         }
209     printf("Enter allocation matrix: ");
210     for(i = 0; i < m; i++)
211         for(j = 0; j < n; j++) {
212             scanf("%d", &s->allocation[i][j]);
213         }
214     printf("\n");
215
216     /* Calculate the need matrix */
217     for(i = 0; i < m; i++)
218         for(j = 0; j < n; j++)
219             s->need[i][j] = s->max[i][j] - s->allocation[i][j];
220
221     /* Calculate the availability vector */
222     for(j = 0; j < n; j++) {
223         int sum = 0;
224         for(i = 0; i < m; i++)
225             sum += s->allocation[i][j];
226         s->available[j] = s->resource[j] - sum;
227     }
228
229     /* Output need matrix and availability vector */
230     printf("Need matrix:\n");
231     for(i = 0; i < m; i++)
232         printf("R%d ", i+1);
233     printf("\n");
234     for(i = 0; i < m; i++) {
235         for(j = 0; j < n; j++)
236             printf("%d ", s->need[i][j]);
237         printf("\n");
238     }
239     printf("Availability vector:\n");
240     for(i = 0; i < n; i++)
241         printf("R%d ", i+1);
242     printf("\n");
243     for(j = 0; j < n; j++)
244         printf("%d ", s->available[j]);
245     printf("\n");
246
247     /* If initial state is unsafe then terminate with error

```

```

248     */
249     int req[n];
250     for(i=0; i<n; i++)
251         req[i]=0;
252     for(i=0; i<m; i++)
253         if (!resource_request(i, req))
254             exit -1;
255
256     /* Seed the random number generator */
257     struct timeval tv;
258     gettimeofday(&tv, NULL);
259     srand(tv.tv_usec);
260
261     /* Create m threads */
262     pthread_t *tid = malloc(m*sizeof(pthread_t));
263     for (i = 0; i < m; i++)
264         pthread_create(&tid[i], NULL, process_thread, (void
265             *) (long) i);
266
267     /* Wait for threads to finish */
268     pthread_exit(0);
269     free(tid);
270
271     /* Free state memory */
272 }
273 // vim: set ts=2 sw=2 et:

```

input.txt

```

1 4
2 3
3
4 9 3 6
5
6 3 2 2
7 6 1 3
8 3 1 4
9 4 2 2
10
11 0 0 0
12 0 0 0
13 0 0 0
14 0 0 0

```

banker/Makefile

```

1 all: banker
2
3 OBJS = banker.o

```

```
4 LIBS= -lpthread
5 CFLAGS= "-ggdb"
6
7 banker: ${OBS}
8         gcc -o $@ ${OBS} ${LIBS}
9
10 clean:
11        rm -rf *o banker
```